






1. Consider the following contour map plotted in 2d.





Note that in the above 2d plot the horizontal axis corresponds to the parameter w and the vertical axis corresponds to the parameter b . Which of the 3d plots below corresponds to the 2d plot shown in the figure above (please see carefully which axis corresponds to w and which to b in the 3d plot).



A.




B.



C.

Solution: Option C is the correct answer.

2. Consider the loss function $\mathcal{L}(w)$ as shown in the figure below. You are interested in finding the minima of this function *i.e.*, the value(s) of w for which the function will take its lowest value. To do so you run gradient descent starting with a random value w_0 (the leftmost red dot in the figure). After running, three steps of gradient descent you have the updated value of w as w_3 . The red dots in the figure show the value of w at each step and the blue dots show the corresponding value of the loss function $\mathcal{L}(w)$. Now, what will happen if you run the 4th step of gradient descent, *i.e.*, if you try to update the value of w using the gradient descent update rule. Assume that the learning rate is 1.



- A. the value of w will increase (*i.e.*, $w_4 > w_3$)
- B. the value of w will remain the same (*i.e.*, $w_4 = w_3$)
- C. the value of w will decrease (*i.e.*, $w_4 < w_3$)

Solution: Option B is the correct answer.

Note that the update rule for gradient descent is given by:

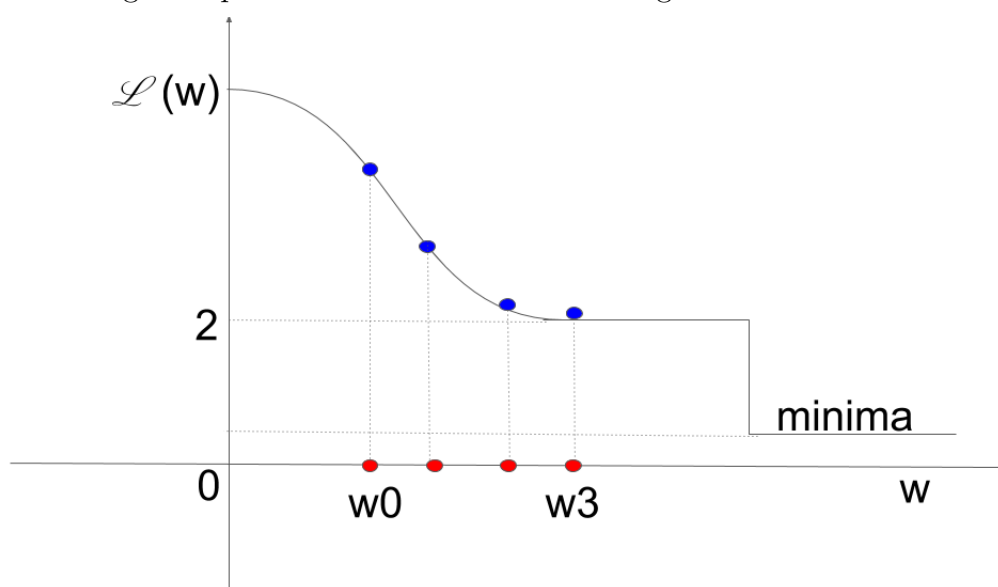
$$w_{t+1} = w_t - \eta \nabla w_t$$

After running 3 steps of gradient descent we have reached a region where the function is absolutely flat, *i.e.*, there is no change in the value of the function in the small neighborhood around w_3 . Hence, the derivative ∇w_t at this point will be zero and the value of w will not get updated. Hence, **Option B** is the correct answer.

3. Continuing the previous question and referring to the same figure again, suppose instead of gradient descent you ran 3 iterations of momentum based gradient descent resulting in the value w_3 as shown in the figure. Note that the update rule of momentum based gradient descent is:

$$\begin{aligned} \text{update}_t &= \gamma \cdot \text{update}_{t-1} + \eta \nabla w_t \\ w_{t+1} &= w_t - \text{update}_t \end{aligned}$$

Assume that the learning rate is 1 and the momentum parameter $\gamma > 0$. Now, what will happen if you run the 4th step of gradient descent, *i.e.*, if you try to update the value of w using the update rule of momentum based gradient descent.



- A. the value of w will increase (*i.e.*, $w_4 > w_3$)
- B. the value of w will remain the same (*i.e.*, $w_4 = w_3$)
- C. the value of w will decrease (*i.e.*, $w_4 < w_3$)

Solution: Option A is the correct answer.

Once again, after running 3 steps of momentum based gradient descent we have reached a region where the function is absolutely flat, *i.e.*, there is no change in the value of the function in the small neighborhood around w_3 . Hence, the derivative ∇w_t at this point will be zero. However, the momentum will still be non-zero and hence w will continue to move in the same direction as it did during the last 3 updates, *i.e.*, the value of w will increase. Hence, **Option A** is the correct answer.

4. Suppose we choose a model $f(x) = \sigma(wx + b)$ which has two parameters w, b . Further, assume that we are trying to learn the parameters of this model using 200 training points. If we use mini-batch gradient descent with a batch size of 10 then how many times will each parameter get updated in one epoch.
- A. 10
 - B. 20
 - C. 100
 - D. 200

Solution: Option B is the correct answer. The parameters will get updated once for every mini-batch and the data is divided into $\frac{200}{10} = 20$ mini-batches. Hence, the parameters will get updated 20 times in one epoch.

5. Note that the update rule for momentum based gradient descent is given by

$$\begin{aligned} \text{update}_t &= \gamma \cdot \text{update}_{t-1} + \eta \nabla w_t \\ w_{t+1} &= w_t - \text{update}_t \end{aligned}$$

Let $\eta = 1$ and $\gamma = 0.9$ and ∇w_1 be the derivative computed at the first time step. If you run momentum based gradient descent for 10 iterations then what fraction of ∇w_1 will be a part of update_{10}

- A. $0.9 \nabla w_1$
- B. $\frac{1}{0.9} \nabla w_1$
- C. $\frac{0.9}{10-1} \nabla w_1$
- D. $(0.9)^{(10-1)} \nabla w_1$

Solution: If we expand the formula for update_t we get

$$\text{update}_t = \gamma \cdot \text{update}_{t-1} + \eta \nabla w_t = \gamma^{t-1} \cdot \eta \nabla w_1 + \gamma^{t-2} \cdot \eta \nabla w_2 + \dots + \eta \nabla w_t$$

Hence, the fraction of ∇w_1 that will be a part of update_t is $\gamma^{t-1} \cdot \eta \nabla w_1$. Substituting, $\gamma = 0.9, \eta = 1$ and $t = 10$ we get $0.9^{(10-1)} \nabla w_1$. Hence, **Option D** is the correct answer.

6. We saw the following update rule for Adam :

$$\begin{aligned}
 m_t &= \beta_1 * m_{t-1} + (1 - \beta_1) * \nabla w_t \\
 v_t &= \beta_2 * v_{t-1} + (1 - \beta_2) * (\nabla w_t)^2 \\
 \hat{m}_t &= \frac{m_t}{1 - \beta_1^t} \quad \hat{v}_t = \frac{v_t}{1 - \beta_2^t} \\
 w_{t+1} &= w_t - \frac{\eta}{\sqrt{\hat{v}_t} + \epsilon} * \hat{m}_t
 \end{aligned}$$

\hat{m}_t, \hat{v}_t are the bias corrected values of m_t, v_t . Suppose, instead of using the above equation for m_t we use the following equation where $0 \leq \alpha_1 \leq 1$ and $0 \leq \beta_1 \leq 1$

$$m_t = \frac{\alpha_1}{\beta_1} * m_{t-1} + \frac{(\beta_1 - \alpha_1)}{\beta_1} * \nabla w_t$$

then what would the bias corrected value of m_t be ?

- A. $\hat{m}_t = \frac{m_t}{\alpha_1^t - \beta_1^t}$
- B. $\hat{m}_t = \frac{\alpha_1^t m_t}{1 - \beta_1^t}$
- C. $\hat{m}_t = \frac{\alpha_1^t m_t}{\alpha_1^t - \beta_1^t}$
- D. $\hat{m}_t = \frac{\beta_1^t m_t}{\beta_1^t - \alpha_1^t}$

Solution: Option D is the correct answer

Note that when,

$$m_t = \beta_1 * m_{t-1} + (1 - \beta_1) * \nabla w_t$$

the bias corrected value was,

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t}$$

Now the new equation for m_t is

$$\begin{aligned}
 m_t &= \frac{\alpha_1}{\beta_1} * m_{t-1} + \frac{(\beta_1 - \alpha_1)}{\beta_1} * \nabla w_t \\
 &= \frac{\alpha_1}{\beta_1} * m_{t-1} + \left(1 - \frac{\alpha_1}{\beta_1}\right) * \nabla w_t
 \end{aligned}$$

in other words β_1 has been replaced by $\frac{\alpha_1}{\beta_1}$. Hence the bias corrected value can also be obtained by replacing β_1 by $\frac{\alpha_1}{\beta_1}$, *i.e.*,

$$\begin{aligned}\hat{m}_t &= \frac{m_t}{1 - \left(\frac{\alpha_1}{\beta_1}\right)^t} \\ &= \frac{\beta_1^t m_t}{\beta_1^t - \alpha_1^t}\end{aligned}$$

7. In this question you will implement the Adam algorithm on toy 2-D dataset which consists of 40 data points, *i.e.*, 40 (x,y) pairs. You can download the dataset using the URL : <https://drive.google.com/file/d/1w6aXg7K7nIv4DBUyWAM-abXvKh2frYRM/view?usp=sharing>

For this question you have to use the squared error loss function which is given as,

$$loss = \frac{1}{2}(\hat{y} - y)^2$$

where \hat{y} is the output of your model given by:

$$\hat{y} = \frac{1}{1 + e^{-(wx+b)}}$$

Now given the following hyperparameter settings,

- learning rate = 0.01
- initial weight, $w = 1$
- initial bias, $b = 1$
- number of iterations = 100
- $\beta_1 = 0.9$
- $\beta_2 = 0.99$

What is the value of the loss at the end of 100 iterations?

- A. loss = 0.058
- B. loss = 0.0
- C. loss = 1.58
- D. loss = 0.58

Solution: Option A is correct. You can find the value of the loss at the end of 100 iterations by using the following code(the one which is given in the lecture video).

```
1 # Importing libraries
2 import pandas as pd
3 import numpy as np
4 import math
5
6 def f(w,b,x):
7     return 1.0 / (1.0 + np.exp(-(w*x + b)))
8
9 def error(w, b): #Calculate loss/error
10    err = 0.0
```

```

11     for x,y in zip(X,Y) :
12         fx = f(w,b,x)
13         err += 0.5 * (fx - y) ** 2
14     return err
15
16 def grad_b(w,b,x,y) :
17     fx = f(w,b,x)
18     return (fx - y) * fx * (1 - fx)
19
20 def grad_w(w,b,x,y) :
21     fx = f(w,b,x)
22     return (fx - y) * fx * (1 - fx) * x
23
24 def run_adam(X,Y,init_w ,init_b ,eta , max_epochs): # ADAM algorithm
25     # Initializations
26     w,b = init_w , init_b
27     m_w, m_b, v_w, v_b, eps, beta1, beta2 = 0,0,0,0,1e-8,0.9,0.999
28
29     for i in range(max_epochs):
30         dw,db = 0,0
31         for x,y in zip(X,Y):
32             dw += grad_w(w,b,x,y)
33             db += grad_b(w,b,x,y)
34         # Compute history
35         m_w = beta1 * m_w + (1-beta1)*dw
36         m_b = beta1 * m_b + (1-beta1)*db
37
38         v_w = beta2 * v_w +(1-beta2)*dw**2
39         v_b = beta2 * v_b +(1-beta2)*db**2
40
41         # Apply bias correction
42         m_w = m_w/(1-math.pow(beta1 , i+1))
43         m_b = m_b/(1-math.pow(beta1 , i+1))
44
45         v_w = v_w/(1-math.pow(beta2 , i+1))
46         v_b = v_b/(1-math.pow(beta2 , i+1))
47
48         # Apply ADAM's update rule
49         w = w - (eta/np.sqrt(v_w + eps)) * m_w
50         b = b - (eta/np.sqrt(v_b + eps)) * m_b
51
52     return w,b
53
54 if __name__ == "__main__":
55     filename = 'A4-Q7_data.csv'
56     df = pd.read_csv(filename) #Loading data
57     X = df['X']
58     Y = df['Y']
59     initial_w = 1
60     initial_b = 1
61     eta = 0.01

```

```

62 max_epochs = 100
63 w,b = run_adam(X, Y, initial_w , initial_b , eta , max_epochs)
64 error = error(w,b)
65 print("error = {}".format(error))

```

The code below is as per the corrected ADAM update equations.(We have updated the lecture slides, please refer to them.)

Solution:

```

1 import pandas as pd
2 import numpy as np
3 #import matplotlib.pyplot as plt
4 import math
5
6 def f(w,b,x):
7     return 1.0 / (1.0 + np.exp(-(w*x + b)))
8
9 def error(w, b): #Calculate loss/error
10    err = 0.0
11    for x,y in zip(X,Y) :
12        fx = f(w,b,x)
13        err += 0.5 * (fx - y) ** 2
14    return err
15
16 def grad_b(w,b,x,y) :
17    fx = f(w,b,x)
18    return (fx - y) * fx * (1 - fx)
19
20 def grad_w(w,b,x,y) :
21    fx = f(w,b,x)
22    return (fx - y) * fx * (1 - fx) * x
23
24 def do_adam(X,Y,init_w ,init_b ,eta , max_epochs):
25    w,b = init_w , init_b
26    m_w, m_b, v_w, v_b, m_w_hat, m_b_hat, v_w_hat, v_b_hat, eps, beta1 ,
27    beta2 = 0,0,0,0,0,0,0,0,1e-8,0.9,0.99
28    for i in range(max_epochs):
29        dw,db = 0,0
30        for x,y in zip(X,Y):
31            dw += grad_w(w,b,x,y)
32            db += grad_b(w,b,x,y)
33            m_w = beta1 * m_w + (1-beta1)*dw
34            m_b = beta1 * m_b + (1-beta1)*db
35
36            v_w = beta2 * v_w +(1-beta2)*dw**2
37            v_b = beta2 * v_b +(1-beta2)*db**2

```

```

38     m_w_hat = m_w/(1-math.pow(beta1,i+1))
39     m_b_hat = m_b/(1-math.pow(beta1,i+1))
40
41     v_w_hat = v_w/(1-math.pow(beta2,i+1))
42     v_b_hat = v_b/(1-math.pow(beta2,i+1))
43
44     w = w - (eta/np.sqrt(v_w_hat + eps)) * m_w_hat
45     b = b - (eta/np.sqrt(v_b_hat + eps)) * m_b_hat
46     return w,b
47
48 if __name__ == "__main__":
49     filename = 'A4-Q8_data.csv'
50     df = pd.read_csv(filename) #Loading data
51     X = df['X']
52     Y = df['Y']
53     initial_w = 1
54     initial_b = 1
55     eta = 0.01
56     max_epochs = 100
57     w,b = do_adam(X, Y, initial_w, initial_b, eta, max_epochs)
58     error = error(w,b)
59     print("error = {}".format(error))

```

On executing this code, the error value which you should get is 0.0036